


Informed search algorithms & Heuristic Functions



Chapter 3
3.5 and 3.6



Outline

- Best-first search
- Greedy best-first search
- A^* search
- Memory bounded Heuristics search
- Learning to search better



Best first search

- Best-first search is an instance of the general TREE-SEARCH or GRAPH-SEARCH algorithm in which a node is selected for expansion based on an **evaluation function**, $f(n)$.
- The evaluation function is interpreted as a cost estimate, so the node with the *lowest* evaluation is expanded first.
- Consider Heuristic function to be arbitrary, nonnegative, problem-specific functions, with one constraint: if n is a goal node, then $h(n)=0$.
- Most best-first algorithms include as a component of f a **heuristic function**, denoted $h(n)$:
 - $h(n)$ = estimated cost of the cheapest path from the state at node n to a goal state.




Best-first search

- Idea: use an **evaluation function** $f(n)$ for each node
 - $f(n)$ provides an estimate for the total cost.
 - Expand the node n with smallest $f(n)$.
- Implementation:

Order the nodes in fringe increasing order of cost.
- Special cases:
 - greedy best-first search
 - A^* search

Romania with straight-line distance






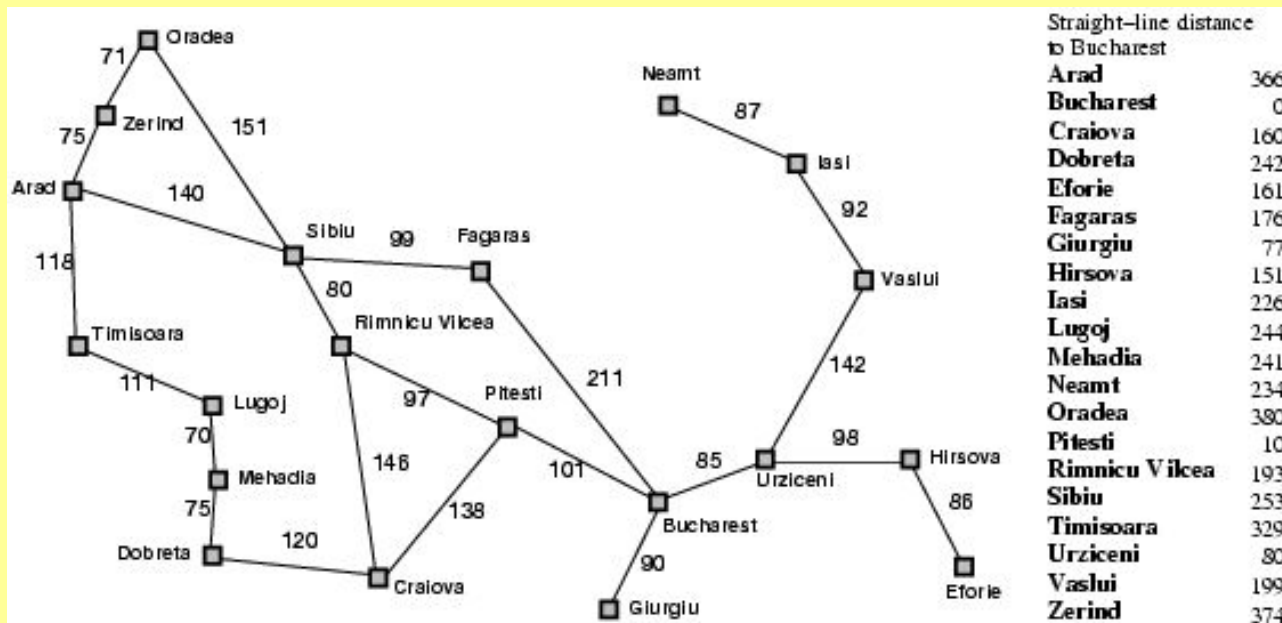
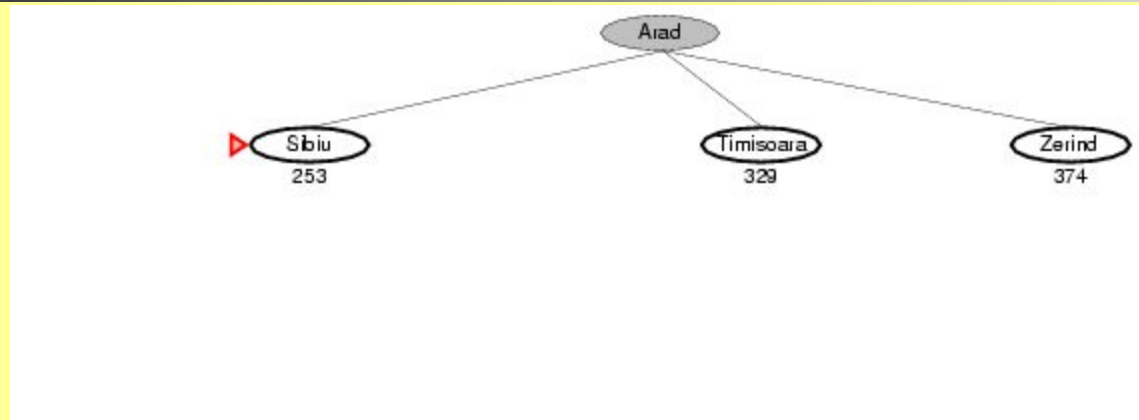
Greedy best-first search

- $f(n)$ = estimate of cost from n to *goal*
- e.g., $f_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal.

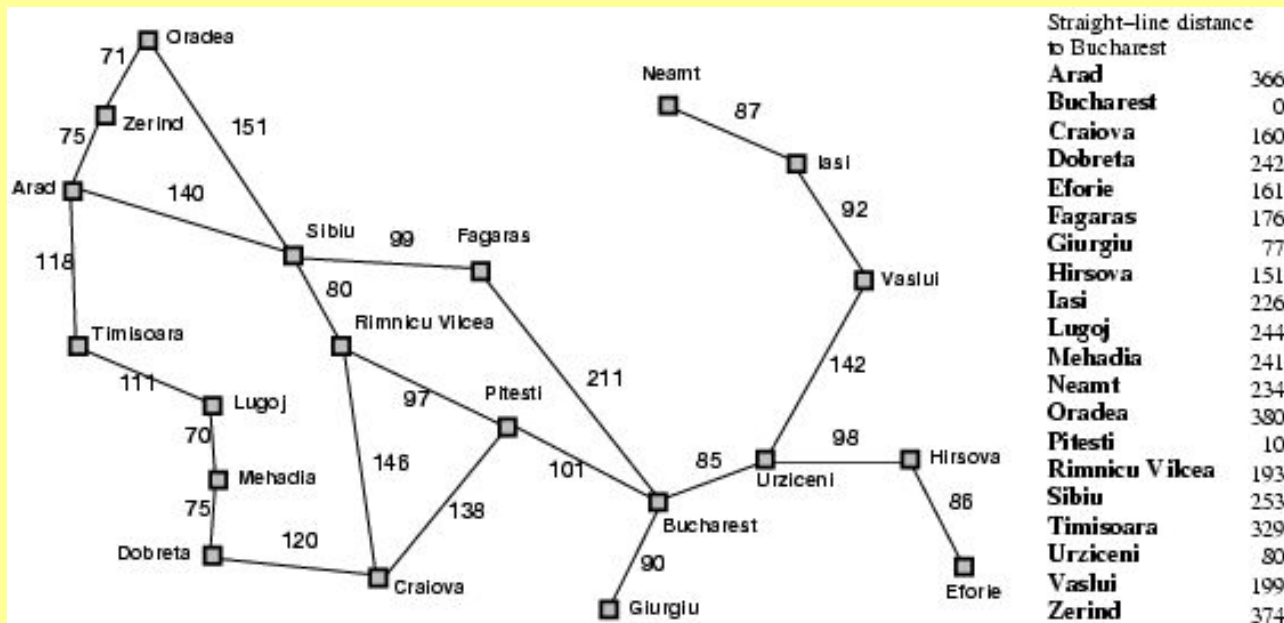

Greedy best-first search example



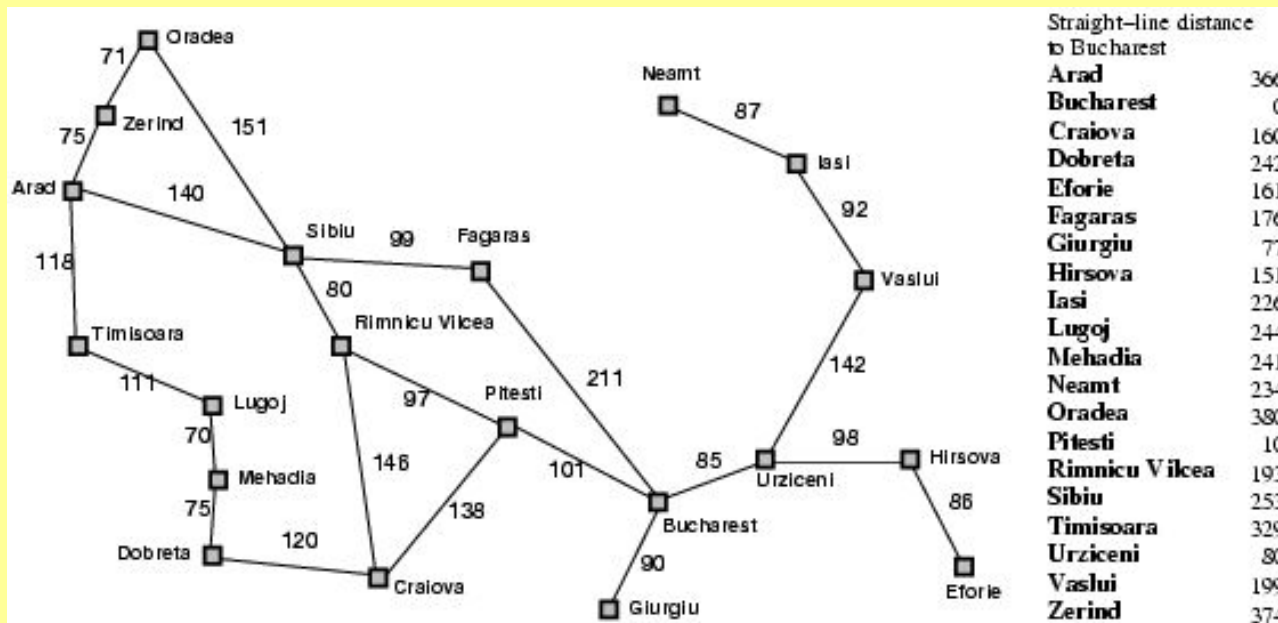
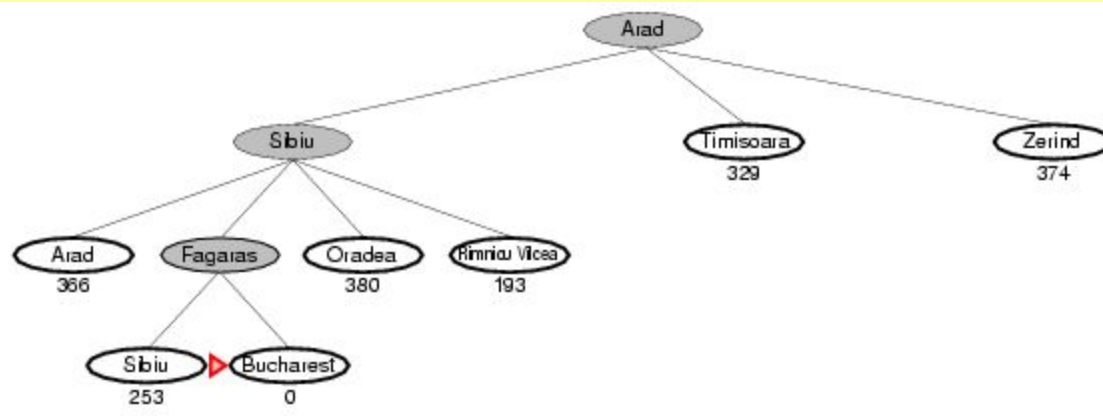
Greedy best-first search example



Greedy best-first search example



Greedy best-first search example



Properties of greedy best-first search

- Complete? No – can get stuck in loops / dead end. (Iasi to Fagarus)
- Time? $O(b^m)$, but a good heuristic can give dramatic improvement
- Space? $O(b^m)$ - keeps all nodes in memory
- Optimal? No
e.g. Arad \square Sibiu \square Rimnicu
Virea \square Pitesti \square Bucharest is shorter!

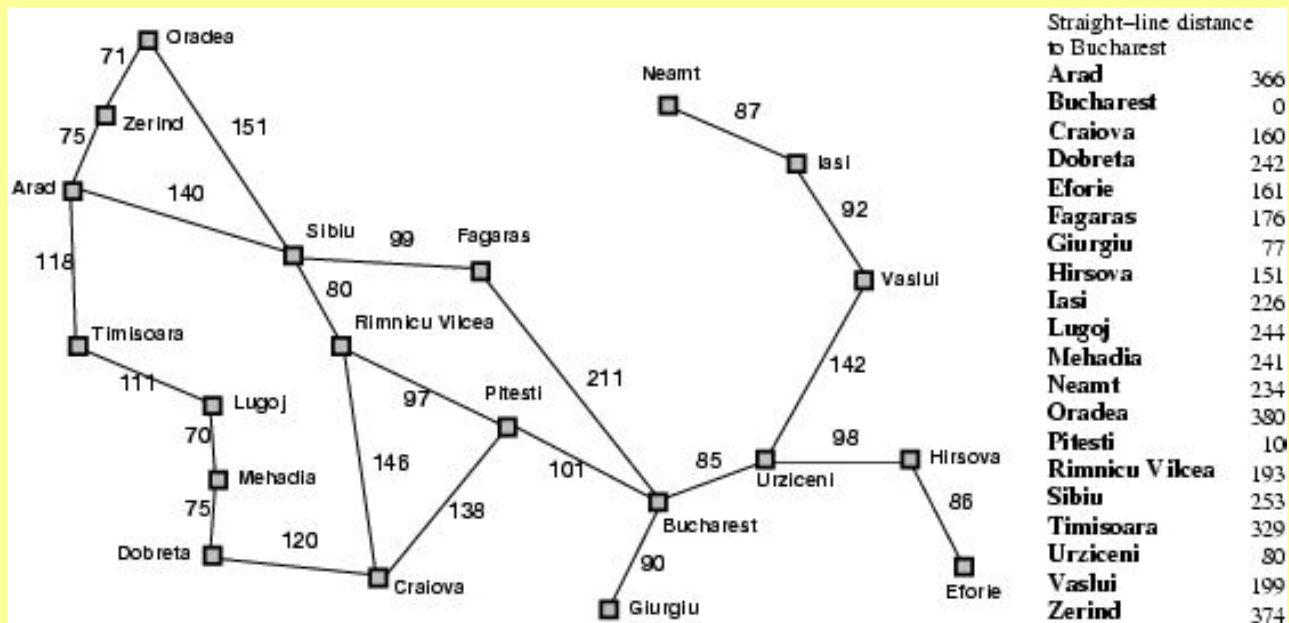


A* search

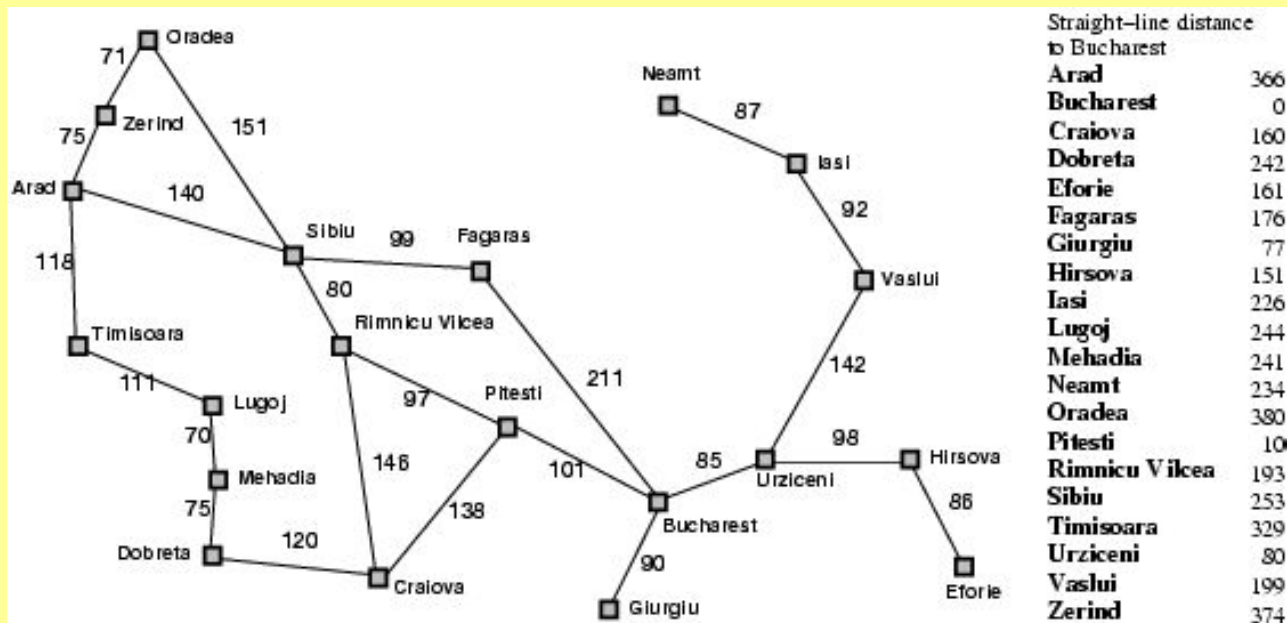

- Idea: avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
- $g(n)$ = cost so far to reach n
- $h(n)$ = estimated cost from n to goal
- $f(n)$ = estimated total cost of path through n to goal
- Best First search has $f(n)=h(n)$

A* search example


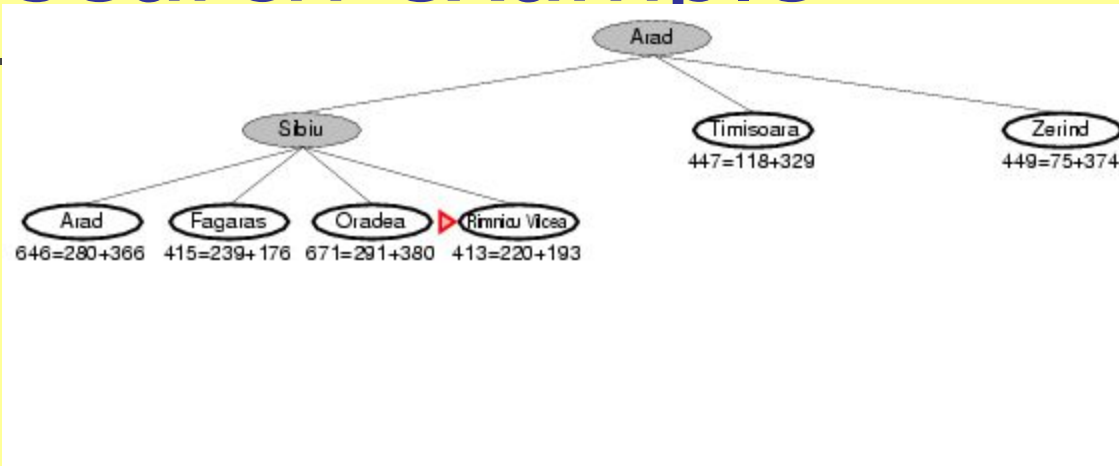
Arad
366=0+366



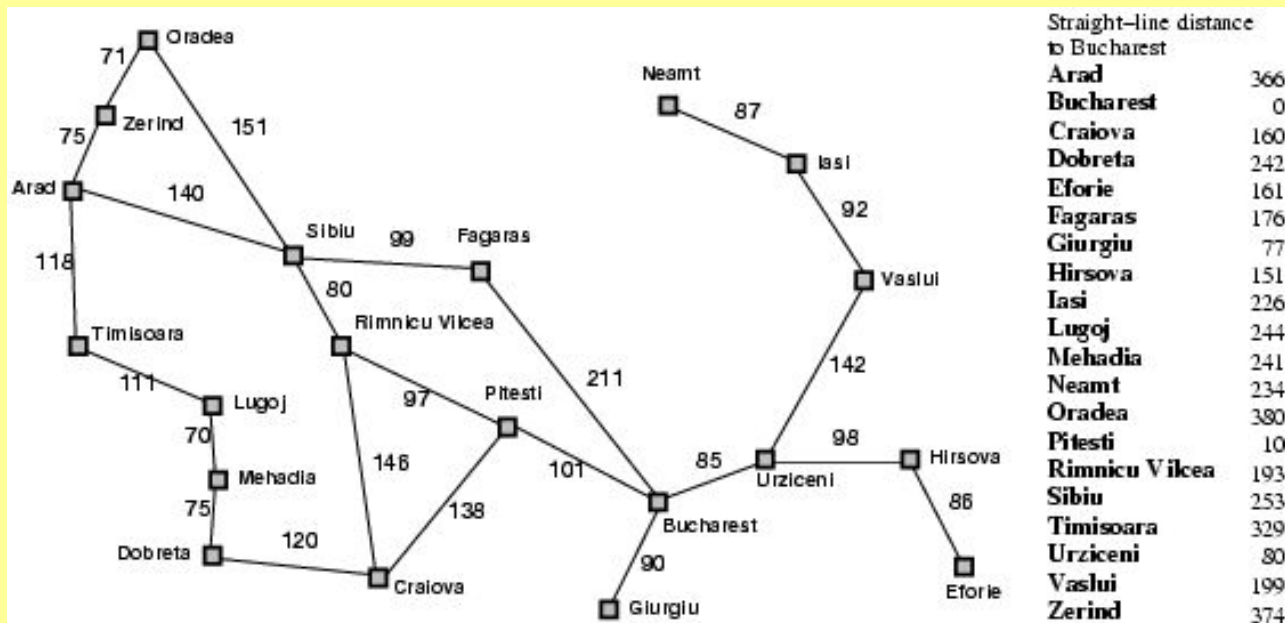
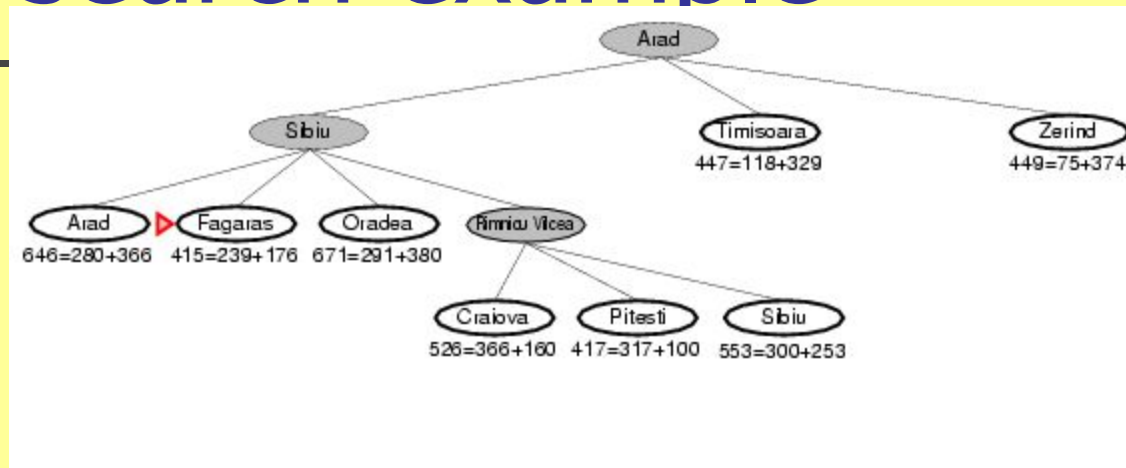
A* search example




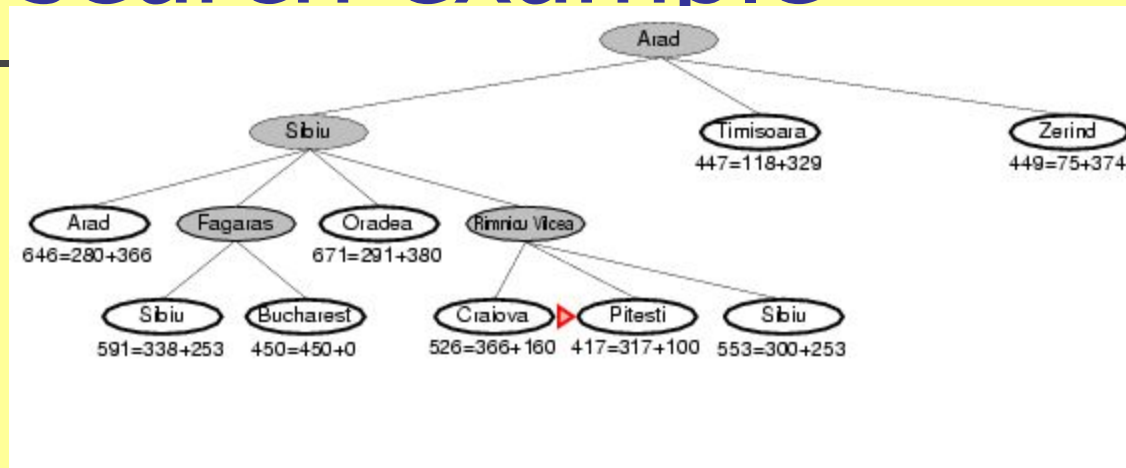
A* search example


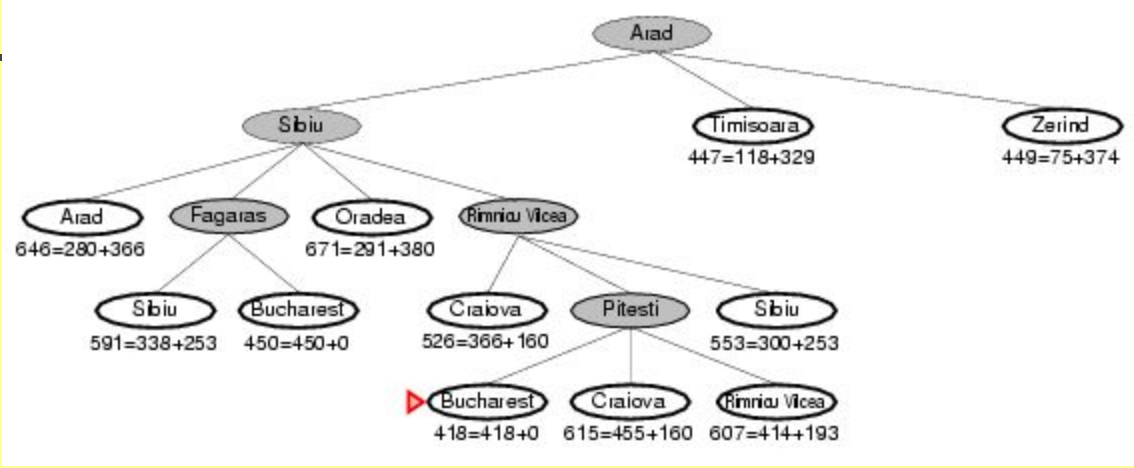


A* search example



A* search example







End of class



Admissible heuristics

- A heuristic $h(n)$ is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n .
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**
- Example: $h_{SLD}(n)$ (never overestimates the actual road distance)
- **Theorem:** If $h(n)$ is admissible, A^* using TREE-SEARCH is **optimal** *while the graph-search version is optimal if $h(n)$ is consistent.*

Consistent heuristics

- A heuristic is **consistent** if for every node n , every successor n' of n generated by any action a ,

$$h(n) \leq c(n,a,n') + h(n')$$

- If h is consistent, we have


$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n,a,n') + h(n') \\ &\geq g(n) + h(n) = f(n) \end{aligned}$$

$$f(n') \geq f(n)$$

- i.e., $f(n)$ is non-decreasing along any path.

- Theorem:**

If $h(n)$ is consistent, A* using GRAPH-SEARCH is optimal




It's the triangle inequality !

keeps all checked nodes
in memory to avoid repeated states

Optimality of A*

- A* expands nodes in order of increasing f value
- Gradually adds " f -contours" of nodes
- Contour i contains all nodes with $f \leq f_i$ where $f_i < f_{i+1}$





Problem of A^*

reach the final state from initial state using A* Algorithm.


Consider $g(n)$ = Depth of node and
 $h(n)$ = Number of misplaced tiles.

| | | |
|---|---|---|
| 2 | 8 | 3 |
| 1 | 6 | 4 |
| 7 | | 5 |

Initial State

| | | |
|---|---|---|
| 1 | 2 | 3 |
| 8 | | 4 |
| 7 | 6 | 5 |

Final State



Initial State

| | | |
|---|---|---|
| 2 | 8 | 3 |
| 1 | 6 | 4 |
| 7 | | 5 |

$g = 0$
 $h = 4$
 $f = 0 + 4 = 4$

| | | |
|---|---|---|
| 2 | 8 | 3 |
| 1 | 6 | 4 |
| | 7 | 5 |

$g = 1$
 $h = 5$
 $f = 1 + 5 = 6$

| | | |
|---|---|---|
| 2 | 8 | 3 |
| 1 | | 4 |
| 7 | 6 | 5 |


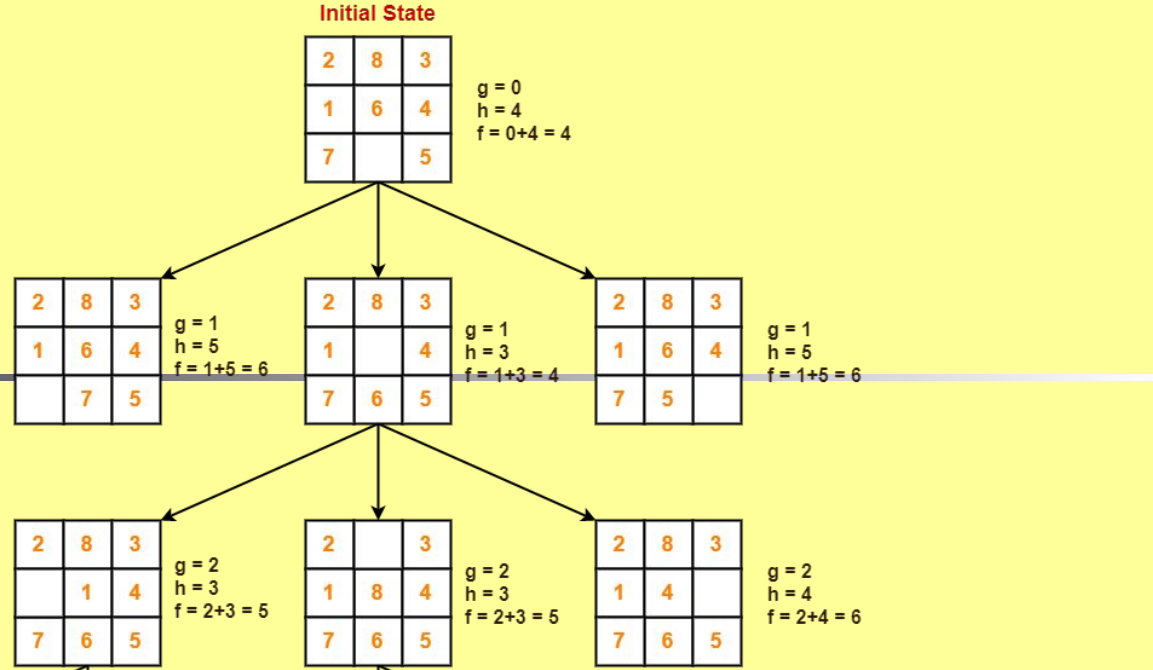

$g = 1$
 $h = 3$
 $f = 1 + 3 = 4$

| | | |
|---|---|---|
| 2 | 8 | 3 |
| 1 | 6 | 4 |
| 7 | 5 | |



$g = 1$
 $h = 5$
 $f = 1 + 5 = 6$







Final State



Final State









Properties of A*

- Complete? Yes (unless there are infinitely many nodes with $f \leq f(G)$, i.e. path-cost $> \epsilon$)
- Time/Space? Exponential b^d
- Optimal? Yes
- Optimally Efficient: Yes (no algorithm with the same heuristic is guaranteed to expand fewer nodes)



Memory Bounded Heuristic Search: Recursive BFS

- How can we solve the memory problem for A^* search?
- Idea: Try something like depth first search, but let's not forget everything about the branches we have partially explored.
- We remember the best f-value we have found so far in the branch we are deleting.



Types of Memory Bounded algorithms

- IDA* - Is practical for many problems with unit step costs and avoids the substantial overhead associated with keeping a sorted queue of nodes. Unfortunately, it suffers from the same difficulties with real valued costs as does the iterative version of uniform-cost search
- RBFS
- MA*

Difference between IDA* and IDDFS



- The main difference between IDA* and standard iterative deepening is that the cutoff used is the f-cost ($g+h$) rather than the depth;
- At each iteration, the cutoff value is the smallest f-cost of any node that exceeded the cutoff on the previous iteration.

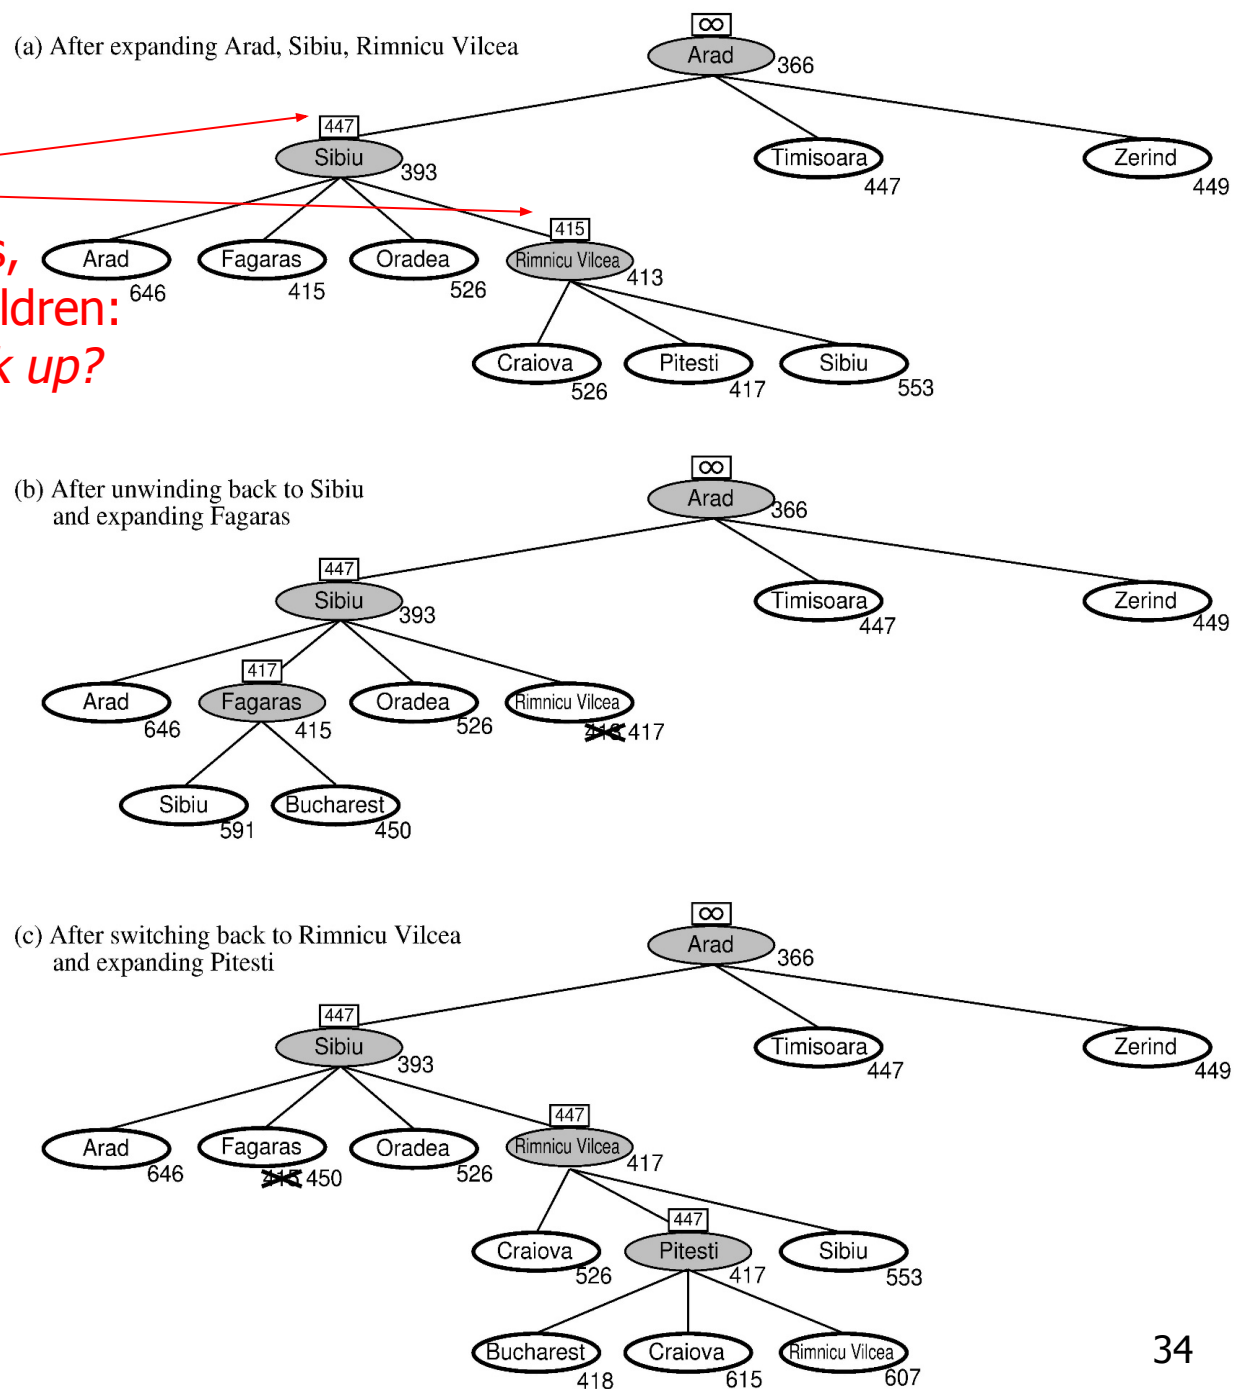
RBFS:

best alternative
over fringe nodes,
which are not children:
do I want to back up?

RBFS changes its mind
very often in practice.

This is because the
 $f=g+h$ become more
accurate (less optimistic)
as we approach the goal.
Hence, higher level nodes
have smaller f -values and
will be explored first.

Problem: We should keep
in memory whatever we can.





End of class



RBFS

- Suffers from excessive node regeneration.
- Like A^* tree search, RBFS is an optimal algorithm if the heuristic function $h(n)$ is admissible.
- Its space complexity is linear in the depth of the deepest optimal solution
- Time complexity is rather difficult to characterize: it depends both on the accuracy of the heuristic function and on how often the best path changes as nodes are expanded.



IDA* and RBFS- Drawback

- IDA* and RBFS suffer from using *too little* memory.
- To utilize available memory
 - **MA*** (memory-bounded A*)
 - **SMA*** (simplified MA*)



Properties of SMA*

- SMA* expands the best leaf and deletes the worst leaf.
- SMA* expands the *newest* best leaf and deletes the *oldest* worst leaf – if F-value same
- SMA* is complete if there is any reachable solution
- It is optimal if any optimal solution is reachable; otherwise, it returns the best reachable solution.
- Time and space complexity is inescapable problem if subset is regenerated.



Learning to search better

- Could an agent *learn* how to search better?
- **Metalevel state space** captures the internal (computational) state of a program that is searching in an **object-level state space** such as Romania.
- For example, the internal state of the A* algorithm consists of the current search tree. Each action in the metalevel state space is a computation step that alters the internal state; for example, each computation step in A* expands a leaf node and adds its successors to the tree.
- **Metalevel learning** algorithm- The goal of learning is to minimize the **total cost** of problem solving, trading off computational expense and path cost.

3.6 Heuristic Functions - Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance/ city block distance (i.e., no. of squares from desired location of each tile)

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

- $h_1(S) = ?$
- $h_2(S) = ?$

Admissible heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance
(i.e., no. of squares from desired location of each tile)

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

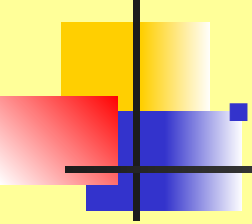
Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State

- $h_1(S) = ?$ 8
- $h_2(S) = ?$ $3+1+2+2+2+3+3+2 = 18$

Dominance

- 
- If $h_2(n) \geq h_1(n)$ for all n (both admissible)
 - then h_2 **dominates** h_1
 - h_2 is better for search: it is guaranteed to expand less nodes.
 - Typical search costs (average number of nodes expanded):
 - $d=12$ IDS = 3,644,035 nodes
 - $A^*(h_1) = 227$ nodes
 - $A^*(h_2) = 73$ nodes
 - $d=24$ IDS = too many nodes
 - $A^*(h_1) = 39,135$ nodes
 - $A^*(h_2) = 1,641$ nodes



Relaxed problems

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution




Generating admissible heuristics from relaxed problems

- If a collection of admissible heuristics $h_1 \dots h_m$ is available for a problem and none of them dominates any of the others, which should we choose?
- As it turns out, we need not make a choice.

$$h(n) = \max\{h_1(n), \dots, h_m(n)\}$$

Generating admissible heuristics from subproblems: Pattern databases



The idea behind **pattern databases** is to store these exact solution costs for every possible subproblem instance—in the example, every possible configuration of the four tiles and the blank.



Learning heuristics from experience

- How could an agent construct to estimate the cost of a solution beginning from the state at node n ?
 - to devise relaxed problems
 - to learn from experience
 - Learning algorithm can be used to construct a function $h(n)$ that can (with luck) predict solution costs for other states that arise during search. - neural nets, decision trees, reinforcement learning, inductive learning methods



End of unit 2

Next class – Logical Agents